

Mini Project: Data Governance Using Unity Catalog - Advanced Capabilities

Objective:

Participants will:

1. Create multiple schemas and tables using Unity Catalog.
2. Set up data governance features like Data Discovery, Data Audit, Data Lineage, and Access Control.
3. Build a secure environment with fine-grained control over data access and visibility.

Task 1: Set Up Unity Catalog Objects with Multiple Schemas

1. Create a Catalog:

Create a catalog named `finance_data_catalog` for storing financial data.

```
CREATE CATALOG finance_data_catalog;
```

2. Create Multiple Schemas:

Create two schemas inside the catalog:

```
transaction_data  
customer_data
```

```
CREATE SCHEMA finance_data_catalog.transaction_data;
```

```
CREATE SCHEMA finance_data_catalog.customer_data;
```

3. Create Tables in Each Schema:

For `transaction_data`, create a table with columns: `TransactionID`, `CustomerID`, `TransactionAmount`, `TransactionDate`.

```
CREATE TABLE  
finance_data_catalog.transaction_data.transactions (  
TransactionID INT,  
  
CustomerID INT,  
  
TransactionAmount DECIMAL(10, 2),  
  
TransactionDate DATE );
```

For `customer_data`, create a table with columns: `CustomerID`, `CustomerName`, `Email`, `Country`.

```
CREATE TABLE finance_data_catalog.customer_data.customers (  
CustomerID INT,  
CustomerName STRING,  
Email STRING,  
Country STRING );
```

Task 2: Data Discovery Across Schemas

1. Explore Metadata:

Search for tables across both schemas and retrieve metadata using SQL commands.

```
SHOW TABLES IN finance_data_catalog.transaction_data;
```

```
SHOW TABLES IN finance_data_catalog.customer_data;
```

2. Data Profiling:

Run SQL queries to perform data profiling on both tables, discovering trends in transaction amounts and customer locations.

```
SELECT AVG(TransactionAmount) AS AverageTransactionAmount  
FROM finance_data_catalog.transaction_data.transactions;
```

```
SELECT Country, COUNT(*) AS CustomerCount FROM  
finance_data_catalog.customer_data.customers GROUP BY Country;
```

3. Tagging Sensitive Data:

Apply tags to sensitive columns such as Email and TransactionAmount for better governance tracking.

```
ALTER TABLE finance_data_catalog.customer_data.customers
```

```
SET TBLPROPERTIES ('tags' = 'sensitive');
```

```
ALTER TABLE finance_data_catalog.transaction_data.transactions
```

```
SET TBLPROPERTIES ('tags' = 'sensitive');
```

Task 3: Implement Data Lineage and Auditing

1. Track Data Lineage:

Merge data from both schemas (transaction_data and customer_data) to generate a comprehensive view.

Use Unity Catalog to trace the data lineage and track changes between these two tables.

2. Audit User Actions:

Enable audit logs for operations performed on the tables and track who accessed or modified the data.

```
CREATE OR REPLACE VIEW finance_data_catalog.merged_data AS  
SELECT
```

```
t.TransactionID,
```

```
t.CustomerID,
```

```
t.TransactionAmount,
```

```
t.TransactionDate,
```

```
c.CustomerName,  
  
c.Email,  
  
c.Country  
  
FROM finance_data_catalog.transaction_data.transactions t  
  
JOIN finance_data_catalog.customer_data.customers c ON  
t.CustomerID = c.CustomerID;
```

Task 4: Access Control and Permissions

1. Set Up Roles and Groups:

Create two groups: DataEngineers and DataAnalysts .

Assign appropriate roles:

DataEngineers should have full access to both schemas and tables. DataAnalysts should have read-only access to the customer_data schema and restricted access to the transaction_data schema.

```
CREATE ROLE DataEngineers;  
CREATE ROLE DataAnalysts;  
GRANT ALL PRIVILEGES ON SCHEMA  
finance_data_catalog.transaction_data TO ROLE DataEngineers;  
GRANT ALL PRIVILEGES ON SCHEMA  
finance_data_catalog.customer_data TO ROLE DataEngineers;  
GRANT SELECT ON SCHEMA  
finance_data_catalog.customer_data TO ROLE DataAnalysts;  
GRANT SELECT ON  
finance_data_catalog.transaction_data.transactions TO ROLE  
DataAnalysts;
```

2. Row-Level Security:

Implement row-level security for the transaction_data schema, allowing only certain users to view high-value transactions.

```
CREATE OR REPLACE VIEW  
finance_data_catalog.transaction_data.filtered_transactions AS  
SELECT * FROM finance_data_catalog.transaction_data.transactions  
WHERE TransactionAmount < 10000;
```

Task 5: Data Governance Best Practices

1. Create Data Quality Rules:

Implement basic data quality rules to ensure that:

Transaction amounts are non-negative.
Customer emails follow the correct format.

```
CREATE OR REPLACE VIEW finance_data_catalog.valid_transactions
AS SELECT * FROM finance_data_catalog.transaction_data.transactions
WHERE TransactionAmount >= 0;
```

```
CREATE OR REPLACE VIEW finance_data_catalog.valid_customers AS
SELECT * FROM finance_data_catalog.customer_data.customers
WHERE Email LIKE '%_@_%._%';
```

2. Validate Data Governance:

Validate all data governance rules by running SQL queries and checking that the lineage and audit logs capture all operations correctly.

```
SELECT * FROM finance_data_catalog.valid_transactions;
```

```
SELECT * FROM finance_data_catalog.valid_customers;
```

Task 6: Data Lifecycle Management

1. Implement Time Travel:

Use Unity Catalog's Delta Time Travel feature to access historical versions of the transaction_data table and restore to a previous state.

```
SELECT * FROM finance_data_catalog.transaction_data.transactions
VERSION AS OF 0;
```

2. Run a Vacuum Operation:

Run a vacuum operation on the tables to clean up old files and ensure the Delta tables are optimized.

```
VACUUM finance_data_catalog.transaction_data.transactions;
VACUUM finance_data_catalog.customer_data.customers;
```

Mini Project: Advanced Data Governance and Security Using Unity Catalog

Objective:

Participants will:

1. Create a multi-tenant data architecture using Unity Catalog.
2. Explore the advanced features of Unity Catalog, including data discovery, data lineage, audit logs, and access control.

Task 1: Set Up Multi-Tenant Data Architecture Using Unity Catalog

1. Create a New Catalog:

Create a catalog named corporate_data_catalog for storing corporate-wide data.

```
CREATE CATALOG corporate_data_catalog;
```

2. Create Schemas for Each Department:

Create three schemas:

```
sales_data
hr_data
finance_data
```

```
CREATE SCHEMA corporate_data_catalog.sales_data;  
CREATE SCHEMA corporate_data_catalog.hr_data;  
CREATE SCHEMA corporate_data_catalog.finance_data;
```

3. Create Tables in Each Schema:

For sales_data : Create a table with columns SalesID , CustomerID , SalesAmount , SalesDate .

```
CREATE TABLE corporate_data_catalog.sales_data.sales (  
SalesID INT,  
  
CustomerID INT,  
  
SalesAmount DECIMAL(10, 2),  
  
SalesDate DATE );
```

For hr_data : Create a table with columns EmployeeID , EmployeeName , Department , Salary .

```
CREATE TABLE corporate_data_catalog.hr_data.employees (  
EmployeeID INT,  
EmployeeName STRING,  
Department STRING,  
Salary DECIMAL(10, 2) );
```

For finance_data : Create a table with columns InvoiceID , VendorID , InvoiceAmount , PaymentDate .

```
CREATE TABLE corporate_data_catalog.finance_data.invoices (  
InvoiceID INT,  
VendorID INT,  
InvoiceAmount DECIMAL(10, 2),  
PaymentDate DATE );
```

Task 2: Enable Data Discovery for Cross-Departmental Data

1. Search for Tables Across Departments:

Use the Unity Catalog interface to search for tables across the sales_data , hr_data , and finance_data schemas.

It will be done using UI.

2. Tag Sensitive Information:

Tag columns that contain sensitive data, such as Salary in the hr_data schema and InvoiceAmount in the finance_data schema.

```
ALTER TABLE corporate_data_catalog.hr_data.employees SET  
TBLPROPERTIES ('tags' = 'sensitive');
```

```
ALTER TABLE corporate_data_catalog.finance_data.invoices SET  
TBLPROPERTIES ('tags' = 'sensitive');
```

3. Data Profiling:

Perform basic data profiling on the tables to analyze trends in sales, employee salaries, and financial transactions.

```
SELECT AVG(SalesAmount) AS AverageSales, COUNT(*) AS  
TotalSales FROM corporate_data_catalog.sales_data.sales;
```

```
SELECT AVG(Salary) AS AverageSalary, COUNT(*) AS  
TotalEmployees FROM corporate_data_catalog.hr_data.employees;
```

```
SELECT AVG(InvoiceAmount) AS AverageInvoice, COUNT(*) AS  
TotalInvoices FROM corporate_data_catalog.finance_data.invoices;
```

Task 3: Implement Data Lineage and Data Auditing

1. Track Data Lineage:

Track data lineage between the sales_data and finance_data schemas by creating a reporting table that merges the sales and finance data. Use Unity Catalog's data lineage feature to visualize how data flows between these tables.

```
CREATE OR REPLACE VIEW  
corporate_data_catalog.sales_finance_report AS SELECT  
  
s.SalesID,  
  
s.CustomerID,  
  
s.SalesAmount,  
  
s.SalesDate,  
  
f.InvoiceID,  
  
f.InvoiceAmount  
  
FROM corporate_data_catalog.sales_data.sales s JOIN  
corporate_data_catalog.finance_data.invoices f ON s.CustomerID =  
f.VendorID;
```

2. Enable Data Audit Logs:

Ensure that all operations (e.g., data reads, writes, and updates) on the hr_data and finance_data tables are captured in audit logs for regulatory compliance.

It will be done using UI.

Task 4: Data Access Control and Security

1. Set Up Roles and Permissions:

Create the following groups:

SalesTeam : Should have access to the sales_data schema only.

FinanceTeam : Should have access to both sales_data and finance_data schemas.

HRTeam : Should have access to the hr_data schema with the ability to update employee records.

```
CREATE ROLE SalesTeam;
```

```
CREATE ROLE FinanceTeam;
```

```
CREATE ROLE HRTeam;
```

```
GRANT USAGE ON SCHEMA corporate_data_catalog.sales_data TO  
ROLE SalesTeam;
```

```
GRANT USAGE ON SCHEMA corporate_data_catalog.finance_data TO  
ROLE FinanceTeam;
```

```
GRANT USAGE ON SCHEMA corporate_data_catalog.hr_data TO ROLE  
HRTeam;
```

```
GRANT UPDATE ON corporate_data_catalog.hr_data.employees TO  
ROLE HRTeam;
```

2. Implement Column-Level Security:

Restrict access to the Salary column in the hr_data schema, allowing only HR managers to view this data.

```
CREATE ROLE HRManager;
```

```
GRANT SELECT (Salary) ON  
corporate_data_catalog.hr_data.employees TO ROLE HRManager;
```

3. Row-Level Security:

Implement row-level security on the sales_data schema to ensure that each sales representative can only access their own sales records.

```
CREATE OR REPLACE VIEW  
corporate_data_catalog.sales_data.filtered_sales AS SELECT *  
FROM corporate_data_catalog.sales_data.sales WHERE SalesRepID  
= CURRENT_USER();
```

Task 5: Data Governance Best Practices

1. Define Data Quality Rules:

Set up data quality rules to ensure:

Sales amounts are positive in the sales_data table.

Employee salaries are greater than zero in the hr_data table. Invoice amounts in the finance_data table match payment records.

```
CREATE OR REPLACE VIEW  
corporate_data_catalog.valid_sales AS SELECT * FROM  
corporate_data_catalog.sales_data.sales WHERE SalesAmount  
> 0;
```

```
CREATE OR REPLACE VIEW  
corporate_data_catalog.valid_employees AS SELECT * FROM  
corporate_data_catalog.hr_data.employees WHERE Salary > 0;
```

```
CREATE OR REPLACE VIEW  
corporate_data_catalog.valid_invoices AS SELECT * FROM  
corporate_data_catalog.finance_data.invoices WHERE  
InvoiceAmount = (SELECT SUM(InvoiceAmount) FROM  
corporate_data_catalog.finance_data.invoices WHERE  
PaymentDate IS NOT NULL);
```

2. Apply Time Travel for Data Auditing:

Use Delta Time Travel to restore the finance_data table to a previous state after an erroneous update and validate the changes using data audit logs.

```
SELECT * FROM corporate_data_catalog.finance_data.invoices  
VERSION AS OF 1;
```

Task 6: Optimize and Clean Up Delta Tables

1. Optimize Delta Tables:

Use the OPTIMIZE command to improve query performance on the sales_data and finance_data tables.

```
OPTIMIZE corporate_data_catalog.sales_data.sales;  
OPTIMIZE corporate_data_catalog.finance_data.invoices;
```

2. Vacuum Delta Tables:

Run a VACUUM operation to remove old and unnecessary data files from the Delta tables, ensuring efficient storage.

```
VACUUM corporate_data_catalog.sales_data.sales;  
  
VACUUM corporate_data_catalog.finance_data.invoices;
```

Mini Project: Building a Secure Data Platform with Unity Catalog

Objective:

Participants will:

1. Set up a secure data platform using Unity Catalog.
2. Explore key data governance features such as Data Discovery, Data Lineage, Access Control, and Audit Logging.

Task 1: Set Up Unity Catalog for Multi-Domain Data Management 1. Create a New

Catalog:

Create a catalog named enterprise_data_catalog to manage data across various domains.

```
CREATE CATALOG enterprise_data_catalog;
```


2. Create Domain-Specific Schemas:

Create the following schemas:

marketing_data
operations_data
it_data

```
CREATE SCHEMA enterprise_data_catalog.marketing_data;  
CREATE SCHEMA enterprise_data_catalog.operations_data;  
CREATE SCHEMA enterprise_data_catalog.it_data;
```

3. Create Tables in Each Schema:

In the marketing_data schema, create a table with columns: CampaignID , CampaignName , Budget , StartDate .

```
CREATE TABLE enterprise_data_catalog.marketing_data.campaigns (  
  CampaignID INT,  
  
  CampaignName STRING,  
  
  Budget DECIMAL(10, 2),  
  
  StartDate DATE );
```

In the operations_data schema, create a table with columns: OrderID , ProductID , Quantity , ShippingStatus .

```
CREATE TABLE enterprise_data_catalog.operations_data.orders (  
  OrderID INT,  
  ProductID INT,  
  Quantity INT,  
  ShippingStatus STRING );
```

In the it_data schema, create a table with columns: IncidentID , ReportedBy , IssueType , ResolutionTime .

```
CREATE TABLE enterprise_data_catalog.it_data.incidents (  
  IncidentID INT,  
  ReportedBy STRING,  
  IssueType STRING,  
  ResolutionTime INT);
```

Task 2: Data Discovery and Classification

1. Search for Data Across Schemas:

Use Unity Catalog's data discovery features to list all tables in the catalog.

Perform a search query to retrieve tables based on data types (e.g., Budget , ResolutionTime).

```
SHOW TABLES IN enterprise_data_catalog.marketing_data;
```

```
SHOW TABLES IN enterprise_data_catalog.operations_data;  
SHOW TABLES IN enterprise_data_catalog.it_data;
```

```
SELECT * FROM  
enterprise_data_catalog.marketing_data.campaigns WHERE  
Budget IS NOT NULL;  
SELECT * FROM enterprise_data_catalog.it_data.incidents WHERE  
ResolutionTime IS NOT NULL;
```

2. Tag Sensitive Information:

Tag the Budget column in marketing_data and ResolutionTime in it_data as sensitive for better data management and compliance.

```
ALTER TABLE  
enterprise_data_catalog.marketing_data.campaigns SET  
TBLPROPERTIES ('tags' = 'sensitive');  
  
ALTER TABLE enterprise_data_catalog.it_data.incidents SET  
TBLPROPERTIES ('tags' = 'sensitive');
```

3. Data Profiling:

Perform basic data profiling to understand trends in marketing budgets and operational shipping statuses.

```
SELECT AVG(Budget) AS AverageBudget, COUNT(*) AS  
TotalCampaigns FROM  
enterprise_data_catalog.marketing_data.campaigns;  
  
SELECT ShippingStatus, COUNT(*) AS StatusCount FROM  
enterprise_data_catalog.operations_data.orders GROUP BY  
ShippingStatus;
```

Task 3: Data Lineage and Auditing

1. Track Data Lineage Across Schemas:

Link the marketing_data with the operations_data by joining campaign performance with product orders.
Use Unity Catalog to track the lineage of the data from marketing campaigns to sales.

```
CREATE OR REPLACE VIEW  
enterprise_data_catalog.marketing_operations_report AS  
SELECT  
m.CampaignID,  
m.CampaignName,  
m.Budget,  
o.OrderID,  
o.ProductID,
```

```
o.Quantity,  
o.ShippingStatus  
FROM enterprise_data_catalog.marketing_data.campaigns m  
JOIN enterprise_data_catalog.operations_data.orders o ON  
m.CampaignID = o.ProductID;
```

2. Enable and Analyze Audit Logs:

Ensure audit logging is enabled to track all operations on tables within the it_data schema. Identify who accessed or modified the data.

It is done using UI.

Task 4: Implement Fine-Grained Access Control

1. Create User Roles and Groups:

Set up the following groups:

MarketingTeam : Access to the marketing_data schema only.

OperationsTeam : Access to both operations_data and marketing_data schemas.

ITSupportTeam : Access to the it_data schema with permission to update issue resolution times.

```
CREATE ROLE MarketingTeam;
```

```
CREATE ROLE OperationsTeam;
```

```
CREATE ROLE ITSupportTeam;
```

```
GRANT USAGE ON SCHEMA  
enterprise_data_catalog.marketing_data TO ROLE  
MarketingTeam;
```

```
GRANT USAGE ON SCHEMA  
enterprise_data_catalog.operations_data TO ROLE  
OperationsTeam;
```

```
GRANT USAGE ON SCHEMA enterprise_data_catalog.it_data  
TO ROLE ITSupportTeam;
```

```
GRANT UPDATE ON enterprise_data_catalog.it_data.incidents  
TO ROLE ITSupportTeam;
```

2. Implement Column-Level Security:

Restrict access to the Budget column in the marketing_data schema, allowing only the MarketingTeam to view it.

```
CREATE ROLE MarketingTEAM;
```

```
GRANT SELECT (Budget) ON  
enterprise_data_catalog.marketing_data.campaigns TO ROLE  
MarketingTEAM;
```

3. Row-Level Security:

Implement row-level security in the operations_data schema to ensure that users from the OperationsTeam can only view orders relevant to their department.

```
CREATE OR REPLACE VIEW
enterprise_data_catalog.operations_data.filtered_orders AS SELECT
* FROM enterprise_data_catalog.operations_data.orders WHERE
DepartmentID = CURRENT_USER();
```

Task 5: Data Governance and Quality Enforcement

1. Set Data Quality Rules:

Define rules for each schema:

marketing_data : Ensure that the campaign budget is greater than zero.

operations_data : Ensure that shipping status is valid (e.g., 'Pending', 'Shipped', 'Delivered').

it_data : Ensure that issue resolution times are recorded correctly and not negative.

```
CREATE OR REPLACE VIEW enterprise_data_catalog.valid_campaigns
AS SELECT * FROM enterprise_data_catalog.marketing_data.campaigns
WHERE Budget > 0;
```

```
CREATE OR REPLACE VIEW enterprise_data_catalog.valid_orders AS
SELECT * FROM enterprise_data_catalog.operations_data.orders
WHERE ShippingStatus IN ('Pending', 'Shipped', 'Delivered');
```

```
CREATE OR REPLACE VIEW enterprise_data_catalog.valid_incidents
AS SELECT * FROM enterprise_data_catalog.it_data.incidents WHERE
ResolutionTime >= 0;
```

2. Apply Delta Lake Time Travel:

Use Delta Lake Time Travel to explore different historical states of the operations_data schema, and revert to an earlier version if required.

```
SELECT * FROM enterprise_data_catalog.operations_data.orders
VERSION AS OF 0;
```

Task 6: Performance Optimization and Data Cleanup

1. Optimize Delta Tables:

Apply OPTIMIZE to the operations_data and it_data schemas to enhance performance for frequent queries.

```
OPTIMIZE enterprise_data_catalog.operations_data.orders;
```

```
OPTIMIZE enterprise_data_catalog.it_data.incidents;
```

2. Vacuum Delta Tables:

Run a VACUUM operation on the Delta tables to clean up old and unnecessary data files.

```
VACUUM enterprise_data_catalog.operations_data.orders;  
VACUUM enterprise_data_catalog.it_data.incidents;
```

Task 1: Raw Data Ingestion

Create a notebook to ingest raw weather data.

The notebook should read a CSV file containing weather data.

Define a schema for the data and ensure that proper data types are used (e.g., City, Date, Temperature, Humidity).

If the raw data file does not exist, handle the error and log it. Save the raw data to a Delta table.

```
from pyspark.sql import SparkSession  
import logging  
import os  
  
logging.basicConfig(level=logging.INFO)  
  
csv_file_path = "/content/weather_data.csv"  
  
if os.path.exists(csv_file_path):  
    try:  
        raw_data = spark.read.csv(csv_file_path, header=True)  
  
        raw_data.write.format("delta").mode("overwrite").save("/delta/weather/raw_data")  
        logging.info("Raw weather data added successfully.")  
    except Exception as e:  
        logging.error(f"Error while adding: {e}")  
else:  
    logging.error("Raw data file does not exist.")
```

Task 2: Data Cleaning

Create a notebook to clean the raw weather data.

Load the data from the Delta table created in Task 1.

Remove any rows that contain missing or null values.

Save the cleaned data to a new Delta table.

```
from pyspark.sql import SparkSession  
import logging  
  
logging.basicConfig(level=logging.INFO)  
  
try:  
    raw_data = spark.read.format("delta").load("/delta/weather/raw_data")  
  
    cleaned_data = raw_data.na.drop()
```

```

cleaned_data.write.format("delta").mode("overwrite").save("/delta/weather/cleaned_data")
logging.info("Cleaned weather data saved successfully.")
except Exception as e:
    logging.error(f"Error during cleaning: {e}")

```

Task 3: Data Transformation

Create a notebook to perform data transformation.
 Load the cleaned data from the Delta table created in Task 2.
 Calculate the average temperature and humidity for each city.
 Save the transformed data to a Delta table.

```

from pyspark.sql.functions import avg

logging.basicConfig(level=logging.INFO)

try:
    cleaned_data = spark.read.format("delta").load("/delta/weather/cleaned_data")

    transformed_data = cleaned_data.groupBy("City").agg(
        avg("Temperature").alias("Average_Temperature"),
        avg("Humidity").alias("Average_Humidity")
    )

    transformed_data.write.format("delta").mode("overwrite").save("/delta/weather/transformed_data")
    logging.info("Transformed weather data saved successfully.")
except Exception as e:
    logging.error(f"Error during transformation: {e}")

```

Task 4: Create a Pipeline to Execute Notebooks

Create a pipeline that sequentially executes the following notebooks: Raw Data Ingestion
 Data Cleaning
 Data Transformation

Handle errors such as missing files or failed steps in the pipeline. Ensure that log messages are generated at each step to track the progress of the pipeline.

By creating separate jobs for the above three tasks pipeline can be created by tasking them sequentially.

Task 1: Raw Data Ingestion

Use the following CSV data to represent daily weather conditions:

```

City,Date,Temperature,Humidity
New York,2024-01-01,30.5,60
Los Angeles,2024-01-01,25.0,65
Chicago,2024-01-01,-5.0,75
Houston,2024-01-01,20.0,80
Phoenix,2024-01-01,15.0,50

```

Load the CSV data into a Delta table in Databricks.

If the file does not exist, handle the missing file scenario and log the error.

```
from pyspark.sql import SparkSession
import logging
import os

logging.basicConfig(level=logging.INFO)

csv_file_path = "/content/weather_data.csv"

if os.path.exists(csv_file_path):
    try:
        # Read the CSV file
        raw_data = spark.read.csv(csv_file_path, header=True, schema=schema)

        raw_data.write.format("delta").mode("overwrite").save("/delta/weather/raw_data")
        logging.info("Raw weather data ingested successfully.")
    except Exception as e:
        logging.error(f"Error during ingestion: {e}")
else:
    logging.error("Raw data file does not exist.")
```

Task 2: Data Cleaning

Create a notebook to clean the ingested weather data.

Handle null or incorrect values in the temperature and humidity columns. After cleaning, save the updated data to a new Delta table.

```
import logging

logging.basicConfig(level=logging.INFO)

try:
    raw_data = spark.read.format("delta").load("/delta/weather/raw_data")

    cleaned_data = raw_data.filter(
        (raw_data.Temperature.isNotNull()) &
        (raw_data.Humidity.isNotNull()) &
        (raw_data.Temperature >= -50) &
        (raw_data.Humidity >= 0) &
        (raw_data.Humidity <= 100)
    )

    cleaned_data.write.format("delta").mode("overwrite").save("/delta/weather/cleaned_data")
```

```
logging.info("Cleaned weather data saved successfully.")
except Exception as e:
    logging.error(f"Error during cleaning: {e}")
```

Task 3: Data Transformation

Transform the cleaned data by calculating the average temperature and humidity for each city.
Save the transformed data into a new Delta table.

```
from pyspark.sql.functions import avg
import logging

logging.basicConfig(level=logging.INFO)

try:
    cleaned_data = spark.read.format("delta").load("/delta/weather/cleaned_data")

    transformed_data = cleaned_data.groupBy("City").agg(
        avg("Temperature").alias("Average_Temperature"),
        avg("Humidity").alias("Average_Humidity")
    )

    transformed_data.write.format("delta").mode("overwrite").save("/delta/weather/trans
formed_data")
    logging.info("Transformed weather data saved successfully.")
except Exception as e:
    logging.error(f"Error during transformation: {e}")
```

Task 4: Build and Run a Pipeline

Create a Databricks pipeline that executes the following notebooks in sequence: Data ingestion (from Task 1)
Data cleaning (from Task 2)
Data transformation (from Task 3)

Ensure each step logs its status and any errors encountered.

Pipeline can be created from joining the above tasks by creating them as jobs.

Task 1: Customer Data Ingestion

Use the following CSV data representing customer transactions:

```
CustomerID,TransactionDate,TransactionAmount,ProductCategory
C001,2024-01-15,250.75,Electronics
C002,2024-01-16,125.50,Groceries
C003,2024-01-17,90.00,Clothing
C004,2024-01-18,300.00,Electronics
C005,2024-01-19,50.00,Groceries
```

Load the CSV data into a Delta table in Databricks.
If the file is not present, add error handling and log an appropriate message.

```
from pyspark.sql.types import StructType, StructField, StringType, DateType, FloatType
import logging
```



```

import os

logging.basicConfig(level=logging.INFO)

csv_file_path = "/dbfs/path/to/customer_data.csv"

if os.path.exists(csv_file_path):

    try:

        customer_data = spark.read.csv(csv_file_path, header=True, schema=schema)

        customer_data.write.format("delta").mode("overwrite").save("/delta/customer/raw_data")

        logging.info("Customer data ingested successfully.")

    except Exception as e:

        logging.error(f"Error during ingestion: {e}")

else:

    logging.error("Customer data file does not exist.")

```

Task 2: Data Cleaning

Create a notebook to clean the ingested customer data.
 Remove any duplicate transactions and handle null values in the TransactionAmount column.
 Save the cleaned data into a new Delta table.

```

import logging

logging.basicConfig(level=logging.INFO)

try:

    raw_data = spark.read.format("delta").load("/delta/customer/raw_data")

    cleaned_data = raw_data.dropDuplicates()

    cleaned_data = cleaned_data.na.drop(subset=["TransactionAmount"])

    cleaned_data.write.format("delta").mode("overwrite").save("/delta/customer/cleaned_data")

    logging.info("Cleaned customer data saved successfully.")

except Exception as e:

    logging.error(f"Error during cleaning: {e}")

```

Task 3: Data Aggregation

Aggregate the cleaned data by ProductCategory to calculate the total transaction amount per category.

Save the aggregated data to a Delta table.

```
from pyspark.sql.functions import sum
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
try:
```

```
    cleaned_data = spark.read.format("delta").load("/delta/customer/cleaned_data")
```

```
    aggregated_data = cleaned_data.groupBy("ProductCategory").agg(
        sum("TransactionAmount").alias("TotalTransactionAmount")
    )
```

```
    aggregated_data.write.format("delta").mode("overwrite").save("/delta/customer/aggregated_data")
```

```
    logging.info("Aggregated customer data saved successfully.")
```

```
except Exception as e:
```

```
    logging.error(f"Error during aggregation: {e}")
```

Task 4: Pipeline Creation

Build a pipeline that:

1. Ingests the raw customer data (from Task 1).
2. Cleans the data (from Task 2).
3. Performs aggregation (from Task 3).

Ensure the pipeline handles missing files or errors during each stage and logs them properly.

Pipeline can be created from joining the above tasks as Jobs.

Task 5: Data Validation

After completing the pipeline, add a data validation step to verify that the total number of transactions matches the sum of individual category transactions.

```
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
try:
```

```
    cleaned_data = spark.read.format("delta").load("/delta/customer/cleaned_data")
```

```
    aggregated_data = spark.read.format("delta").load("/delta/customer/aggregated_data")
```

```
    total_transactions = cleaned_data.count()
```

```
    total_aggregated =
    aggregated_data.agg(sum("TotalTransactionAmount")).collect()[0][0]
```

```
    if total_transactions == total_aggregated:
```

```
        logging.info("Validation successful: Total transactions match.")
```

```
    else:
```

```
        logging.error("Validation failed: Total transactions do not match.")
```

```
except Exception as e:
```

```
    logging.error(f"Error during validation: {e}")
```

Task 1: Product Inventory Data Ingestion

Use the following CSV data to represent product inventory information:

```
ProductID,ProductName,StockQuantity,Price,LastRestocked
P001,Laptop,50,1500.00,2024-02-01
P002,Smartphone,200,800.00,2024-02-02
P003,Headphones,300,100.00,2024-01-29
P004,Tablet,150,600.00,2024-01-30
P005,Smartwatch,100,250.00,2024-02-03
```

Load this CSV data into a Delta table in Databricks.

Handle scenarios where the file is missing or corrupted and log the error accordingly.

```
from pyspark.sql import SparkSession
import logging
import os

logging.basicConfig(level=logging.INFO)

csv_file_path = "/dbfs/path/to/product_inventory.csv"

if os.path.exists(csv_file_path):
    try:
        product_data = spark.read.csv(csv_file_path, header=True,
        schema=schema)

        product_data.write.format("delta").mode("overwrite").save("/delta/inventory
        /raw_data")

        logging.info("Product inventory data ingested successfully.")
    except Exception as e:
        logging.error(f"Error during ingestion: {e}")
    else:
        logging.error("Product inventory data file does not exist.")
```

Task 2: Data Cleaning

Clean the ingested product data:

- Ensure no null values in StockQuantity and Price columns.
- Remove any records with StockQuantity less than 0.
- Save the cleaned data to a new Delta table.

```
import logging

logging.basicConfig(level=logging.INFO)

try:
    raw_data = spark.read.format("delta").load("/delta/inventory/raw_data")

    cleaned_data = raw_data.na.drop(subset=["StockQuantity", "Price"])
```

```
cleaned_data = cleaned_data.filter(cleaned_data.StockQuantity >= 0)
```

```
cleaned_data.write.format("delta").mode("overwrite").save("/delta/inventory/cleaned_data")
```

```
logging.info("Cleaned product inventory data saved successfully.")
```

```
except Exception as e:
```

```
logging.error(f"Error during cleaning: {e}")
```

Task 3: Inventory Analysis

Create a notebook to analyze the inventory data:

Calculate the total stock value for each product (StockQuantity * Price).

Find products that need restocking (e.g., products with StockQuantity < 100).

Save the analysis results to a Delta table.

```
from pyspark.sql.functions import col, expr
```

```
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
try:
```

```
cleaned_data = spark.read.format("delta").load("/delta/inventory/cleaned_data")
```

```
analysis_data = cleaned_data.withColumn(
```

```
"TotalStockValue", col("StockQuantity") * col("Price")
```

```
)
```

```
products_needing_restock = analysis_data.filter(col("StockQuantity") < 100)
```

```
analysis_data.write.format("delta").mode("overwrite").save("/delta/inventory/analysis_data")
```

```
products_needing_restock.write.format("delta").mode("overwrite").save("/delta/inventory/products_needing_restock")
```

```
logging.info("Inventory analysis completed and results saved successfully.")
```

```
except Exception as e:
```

```
logging.error(f"Error during analysis: {e}")
```

Task 4: Build an Inventory Pipeline

Build a Databricks pipeline that:

1. Ingests the product inventory data (from Task 1).
2. Cleans the data (from Task 2).
3. Performs inventory analysis (from Task 3).

Ensure the pipeline logs errors if any step fails and handles unexpected issues such as missing data.

Pipelines can be created from joining the above tasks as jobs.

Task 5: Inventory Monitoring

Create a monitoring notebook that checks the Delta table for any products that need restocking (e.g., StockQuantity < 50).

The notebook should send an alert if any product is below the threshold.

```
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
try:
```

```
    products_needing_restock =  
    spark.read.format("delta").load("/delta/inventory/products_needing_restock")
```

```
    if products_needing_restock.count() > 0:
```

```
        logging.warning("Attention: Some products need restocking!")
```

```
        products_needing_restock.show() # Display products needing restock
```

```
    else:
```

```
        logging.info("All products are adequately stocked.")
```

```
except Exception as e:
```

```
    logging.error(f"Error during monitoring: {e}")
```

Task 1: Employee Attendance Data Ingestion

Use the following CSV data representing employee attendance logs:

```
EmployeeID,Date,CheckInTime,CheckOutTime,HoursWorked  
E001,2024-03-01,09:00,17:00,8  
E002,2024-03-01,09:15,18:00,8.75  
E003,2024-03-01,08:45,17:15,8.5  
E004,2024-03-01,10:00,16:30,6.5  
E005,2024-03-01,09:30,18:15,8.75
```

Ingest this CSV data into a Delta table in Databricks.

Handle potential issues, such as the file being missing or having inconsistent columns, and log these errors.

```
from pyspark.sql import SparkSession
import logging
import os

logging.basicConfig(level=logging.INFO)

if os.path.exists(csv_file_path):
    try:
        attendance_data = spark.read.csv(csv_file_path, header=True,
schema=schema)

        attendance_data.write.format("delta").mode("overwrite").save("/delta/attendance/
raw_data")

        logging.info("Employee attendance data ingested successfully.")
    except Exception as e:
        logging.error(f"Error during ingestion: {e}")
    else:
        logging.error("Employee attendance data file does not exist.")
```

Task 2: Data Cleaning

Clean the ingested attendance data:

- Remove any rows with null or invalid values in the CheckInTime or CheckOutTime columns.
- Ensure the HoursWorked column is calculated correctly (CheckOutTime - CheckInTime).
- Save the cleaned data into a new Delta table.

```
from pyspark.sql.functions import col, to_timestamp

logging.basicConfig(level=logging.INFO)

try:
    raw_data = spark.read.format("delta").load("/delta/attendance/raw_data")

    cleaned_data = raw_data.na.drop(subset=["CheckInTime", "CheckOutTime"])

    cleaned_data = cleaned_data.withColumn(
        "CheckInTime", to_timestamp(col("Date") + " " + col("CheckInTime"),
"yyyy-MM-dd HH:mm")
    ).withColumn(
        "CheckOutTime", to_timestamp(col("Date") + " " + col("CheckOutTime"),
"yyyy-MM-dd HH:mm")
    ).withColumn(
        "HoursWorked", (col("CheckOutTime").cast("long") -
col("CheckInTime").cast("long")) / 3600
    )

    cleaned_data.write.format("delta").mode("overwrite").save("/delta/attendance/cl
eaned_data")
```

```
logging.info("Cleaned employee attendance data saved successfully.")
except Exception as e:
    logging.error(f"Error during cleaning: {e}")
```

Task 3: Attendance Summary

Create a notebook that summarizes employee attendance:

Calculate the total hours worked by each employee for the current month. Find employees who have worked overtime (e.g., more than 8 hours on any given day).
Save the summary to a new Delta table.

```
from pyspark.sql.functions import *
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
try:
    cleaned_data =
    spark.read.format("delta").load("/delta/attendance/cleaned_data")
```

```
    summary_data = cleaned_data.groupBy("EmployeeID").agg(
        sum("HoursWorked").alias("TotalHoursWorked"),
        sum(when(col("HoursWorked") > 8, 1).otherwise(0)).alias("OvertimeDays")
    )
```

```
summary_data.write.format("delta").mode("overwrite").save("/delta/attendance/summary_data")
logging.info("Employee attendance summary calculated and saved successfully.")
except Exception as e:
    logging.error(f"Error during summarizing: {e}")
```

Task 4: Create an Attendance Pipeline

Build a pipeline in Databricks that:

1. Ingests employee attendance data (from Task 1).
2. Cleans the data (from Task 2).
3. Summarizes the attendance and calculates overtime (from Task 3). Ensure the pipeline logs

errors and handles scenarios like missing data.

Pipeline can be created by joining the above three tasks as jobs.

Task 5: Time Travel with Delta Lake

Implement time travel using Delta Lake:

Roll back the attendance data to a previous version (e.g., the day before a change was made).

Use the DESCRIBE HISTORY command to inspect the changes made to the Delta table.

```
from pyspark.sql import SparkSession
import logging
```

```
logging.basicConfig(level=logging.INFO)
```

```
version_number = 0
```

```
try:
    previous_version_data =
spark.read.format("delta").option("versionAsOf",
version_number).load("/delta/attendance/cleaned_data")

previous_version_data.write.format("delta").mode("overwrite").save("/
delta/attendance/cleaned_data_rollback")

    logging.info(f"Successfully rolled back to version
{version_number}.")

    history = spark.sql("DESCRIBE HISTORY
delta.`/delta/attendance/cleaned_data`")
    history.show(truncate=False)
except Exception as e:
    logging.error(f"Error during time travel: {e}")
```